



Image created using DALL-E

# SPACESHIP TITANIC

Melek Derman

12/09/2024



# What is Spaceship Titanic dataset?

- Spaceship Titanic is a Kaggle competition to help learners understand the basics of Machine Learning.
- It is a classification problem.
- It has more than 8000 data in train set.



# Problem Description

- It is the year 2912. A transmission from 4 lightyears away has been received.
- An interstellar passenger liner, the Spaceship Titanic, was carrying passengers from our solar system to three newly habitable exoplanets orbiting nearby stars.
- Devastatingly, the ship collided with a spacetime anomaly hidden within a dust cloud, resulting in half of the passengers being transported to an alternate dimension!
- Our goal is to retrieve the lost passengers by predicting which passengers were transported using the records. Help save them and change history!



Image created using DALL-E



# Dataset and Features

- There are 8,693 passenger records
- 6 numerical and 6 categorical features
- Numerical Features:
  - Age, RoomService, FoodCourt, ShoppingMall, Spa, VRDeck
- Categorical Features:
  - HomePlanet, CryoSleep, Cabin, Destination, VIP, Name
- The features are mostly unbalanced.
- There are missing values in all categories.

Train data info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 8693 entries, 0 to 8692

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	8693 non-null	object
1	HomePlanet	8492 non-null	object
2	CryoSleep	8476 non-null	object
3	Cabin	8494 non-null	object
4	Destination	8511 non-null	object
5	Age	8514 non-null	float64
6	VIP	8490 non-null	object
7	RoomService	8512 non-null	float64
8	FoodCourt	8510 non-null	float64
9	ShoppingMall	8485 non-null	float64
10	Spa	8510 non-null	float64
11	VRDeck	8505 non-null	float64
12	Name	8493 non-null	object

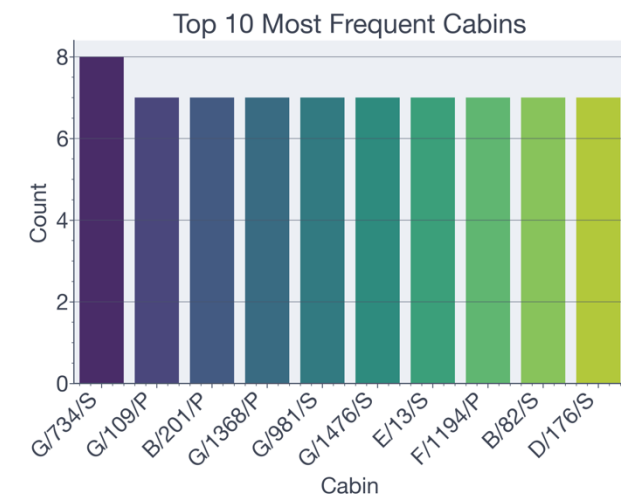
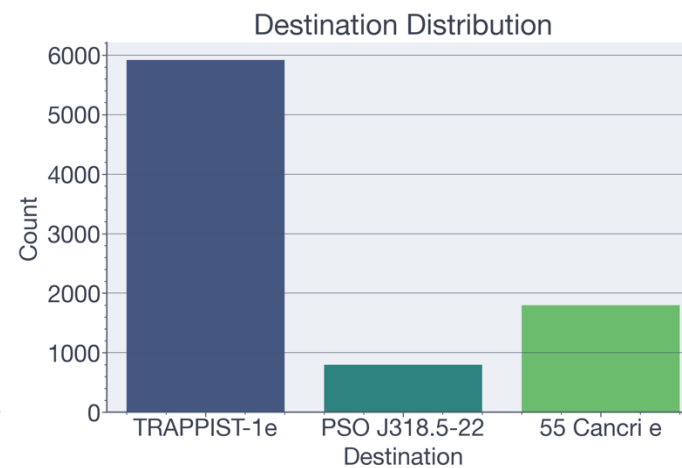

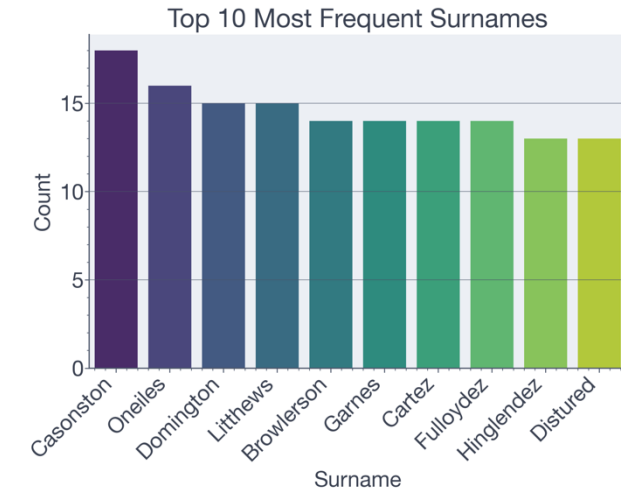


Number of unique values in each column:

PassengerId	8693	Missing Values:
-------------	------	-----------------

HomePlanet	3	HomePlanet	201
CryoSleep	2	CryoSleep	217
Cabin	6560	Cabin	199
Destination	3	Destination	182
Age	80	Age	179
VIP	2	VIP	203
RoomService	1273	RoomService	181
FoodCourt	1507	FoodCourt	183
ShoppingMall	1115	ShoppingMall	208
Spa	1327	Spa	183
VRDeck	1306	VRDeck	188
Name	8473	Name	200
Transported	2		




# - Categorical Features



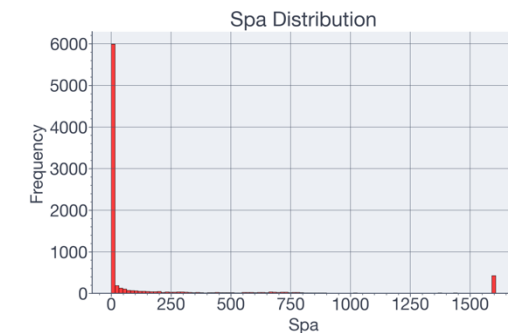
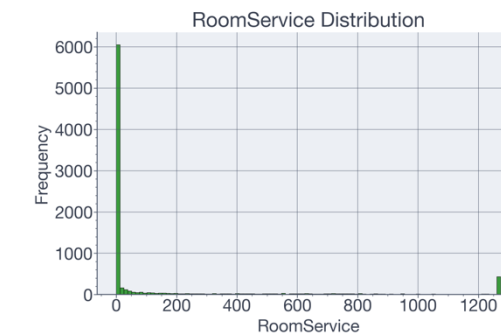
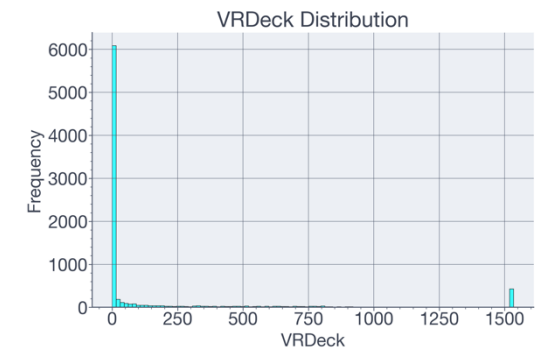
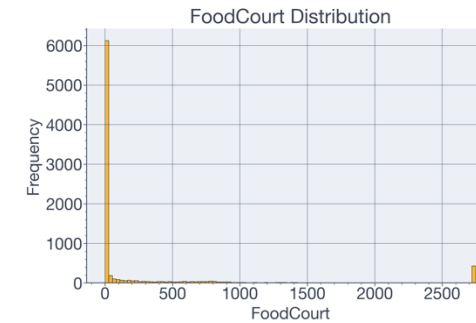
\* CryoSleep indicates whether the passenger elected to be put into suspended animation for the duration of the voyage.



# - Numerical Features



\* RoomService, FoodCourt, ShoppingMall, Spa, VRDeck - Amount the passenger has billed at each of the *Spaceship Titanic*'s many luxury amenities.





# Methodology for Features

- The “Name” column was split into individual words, enabling the analysis of family members.
- For missing values, the ‘mean’ strategy was used for numerical columns, and the ‘most\_frequent’ strategy was applied for categorical features (Simple Imputer)
- Numerical Features → Standard Scaler
- Categorical Features → OneHotEncoder

```
numerical_features = X_train.select_dtypes(include=['float64']).columns.tolist()
categorical_features = X_train.select_dtypes(include=['object']).columns.tolist()

numerical_features = [col for col in num_cols if col in X.columns]
categorical_features = [col for col in X.columns if col not in numerical_features]

numerical_transformer = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
])

preprocessor = ColumnTransformer(transformers=[
    ('num', numerical_transformer, numerical_features),
    ('cat', categorical_transformer, categorical_features)
])
```



# Methodology – k Nearest Neighbor Classifier


- The nearest neighbor classifier is a simple machine learning algorithm that is often used for classification tasks.
- It works by finding the "nearest" training example to a new, unlabeled example in the feature space, and assigning the label of that nearest example to the new example. Predictions are based on the similarity of a new data point to the data points in the training set.
- The algorithm is called "lazy" because it does not actually learn a model from the training data; instead, it simply memorizes the training examples and uses them to make predictions at test time.





# Methodology \_ Nearest Neighbor Classifier

- For a given data point (test point), KNN calculates its distance to all other points in the dataset.
- It selects the k nearest neighbors (based on the smallest distances).
- KNN assumes that similar data points are close to each other in the feature space. Therefore, the prediction for a new data point is based on its proximity to other known data points.





# Methodology K-NN Hyperparameters


- **Number of Neighbors (k):** Determines how many closest neighbors are considered when making a prediction.
  - Small k: The model is sensitive to noise and may overfit.
  - Large k: The model becomes too generalized and may underfit.
- **Distance Metric:** The choice of distance metric determines how the algorithm measures similarity between points.

- **Euclidean Distance:** Works well when the features have similar scales.
- **Manhattan Distance:** More robust when dealing with grid-like data or features with varying scales

Now the Euclidean distance is simply the  $\ell_2$ -norm:      and the Manhattan distance is simply the  $\ell_1$ -norm:

$$\ell_2(x, y) = \sqrt{|x_1 - y_1|^2 + \dots + |x_d - y_d|^2} \quad \ell_1(x, y) = |x_1 - y_1| + \dots + |x_d - y_d|$$

- **Weighting of Neighbors:** Neighbors can be weighted based on their distance to the test point:
  - Uniform Weighting: All neighbors have equal importance.
  - Distance Weighting: Closer neighbors have more influence.



Manhattan vs. Euclidean Distances



# Methodology - Tuning

To optimize the performance of the KNN algorithm, we need to tune the hyperparameters:

## **Split the Data:**

- Use the training data to tune the hyperparameters using cross-validation and evaluate the performance on a separate test set.

## **Hyperparameter Tuning:**

- k: Test odd values such as [1, 3, 5, 7, 9, 11, ...] to avoid tie-breaking votes.
- Distance metric: Try Euclidean, Manhattan.
- Weighting: Test both uniform and distance-based weighting.
- Use cross-validation to test each combination of hyperparameters and select the combination that gives the best performance.

## **Optimal Hyperparameters:**

- After testing various combinations, select the hyperparameters that maximize a performance metric (e.g., accuracy, precision, recall).


## **Final Evaluation:**

- Train the KNN model on the entire training set using the optimal hyperparameters.
- Evaluate the performance on the test set and submit the test predictions to Kaggle.




# Results

Training Set Confusion Matrix for k-nn



Test Set Confusion Matrix for k-nn





# Results

Best hyperparameters: n\_neighbors=41, metric=manhattan, weight=uniform with a balanced accuracy of 0.7828  
Balanced Accuracy for k-NN on train set: 0.7953805283288025  
Accuracy for k-NN on train set: 0.7955013736263736  
Precision for k-NN on train set: 0.7957872343239749  
Recall for k-NN on train set: 0.7953805283288025  
Balanced Accuracy for k-NN on test set: 0.786617695657245  
Accuracy for k-NN on test set: 0.7866852561868247  
Precision for k-NN on test set: 0.7867495261905526  
Recall for k-NN on test set: 0.786617695657245

Cross-validation balanced accuracy scores for each combination:

n_neighbors=1, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7109	n_neighbors=41, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7813
n_neighbors=3, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7476	n_neighbors=43, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7820
n_neighbors=5, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7622	n_neighbors=45, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7808
n_neighbors=7, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7694	n_neighbors=47, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7792
n_neighbors=9, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7689	n_neighbors=49, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7796
n_neighbors=11, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7699	n_neighbors=51, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7798
n_neighbors=13, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7706	n_neighbors=53, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7812
n_neighbors=15, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7785	n_neighbors=55, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7806
n_neighbors=17, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7784	n_neighbors=57, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7807
n_neighbors=19, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7771	n_neighbors=59, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7808
n_neighbors=21, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7772	n_neighbors=61, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7803
n_neighbors=23, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7778	n_neighbors=63, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7796
n_neighbors=25, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7783	n_neighbors=65, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7807
n_neighbors=27, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7788	n_neighbors=67, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7807
n_neighbors=29, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7797	n_neighbors=69, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7812
n_neighbors=31, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7800	n_neighbors=71, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7817
n_neighbors=33, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7805	n_neighbors=73, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7814
n_neighbors=35, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7788	n_neighbors=75, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7800
n_neighbors=37, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7819	n_neighbors=77, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7806
n_neighbors=39, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7805	n_neighbors=79, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7801
n_neighbors=41, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7793	n_neighbors=81, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7808
n_neighbors=43, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7792	n_neighbors=83, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7799
n_neighbors=45, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7821	n_neighbors=85, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7803
n_neighbors=47, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7823	n_neighbors=87, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7808
n_neighbors=49, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7816	n_neighbors=89, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7796
n_neighbors=51, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7809	n_neighbors=91, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7800
n_neighbors=53, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7795	n_neighbors=93, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7803
n_neighbors=55, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7799	n_neighbors=95, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7801
n_neighbors=57, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7811	n_neighbors=97, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7803
n_neighbors=59, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7809	n_neighbors=99, metric=manhattan, weight=distance: Mean Balanced Accuracy = 0.7802
n_neighbors=61, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7808	
n_neighbors=63, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7797	
n_neighbors=65, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7806	
n_neighbors=67, metric=euclidean, weight=uniform: Mean Balanced Accuracy = 0.7816	





Oregon State University  
College of Engineering



# Thank you!

# Questions?

Image created using DALL-E



# Kaggle Test Results



**final\_submission\_svm.csv**

Complete · 7h ago

**0.79331**



**final\_submission.csv**

Complete · 8h ago · knn

**0.75941**



**sample\_submission.csv**

Complete · 5d ago

**0.49310**